

Where Things Stand

Use an AI coding agent to autonomously convert ROCKE-3D (NASA GISS's Fortran GCM) to Python/JAX, and measure accuracy, CPU/GPU performance, conversion cost, and gotchas — validated against real production restart data (P2SAoM40).

14 / 17

Modules faithfully ported

Validated directly against real Fortran. 3 remain placeholders — see slide 3.

1e-9–1e-3

Accuracy vs. real Fortran

PBL + DRYCNV, same inputs, CPU and GPU — floating-point-level agreement.

2.3–6.3×

GPU speedup vs. real Fortran

Measured, not estimated — first real GPU run, 2026-09-20 (kernel-level).

Mixed on CPU

JAX is not a blanket CPU win

Fortran beats JAX-CPU on DRYCNV. The case for JAX here is GPU, not CPU.

The comparison that matters:

Every number here is **JAX vs. the real production Fortran model** — not JAX vs. a second Python reimplementation. NumPy-relative comparisons exist elsewhere in this project's history but are dropped from this summary as answering the wrong question.

PERFORMANCE


Real Fortran vs. JAX — CPU and GPU

Kernel-level, measured (mean of 100 calls)

DRYCNV



PBL similarity



Same shared inputs on every leg. Grid: P2SAoM40's real 72x46x40.

Full physics chain (CPU only)

PBL + radiation + surface + ground, driven by P2SAoM40's real restart state, chained in the real per-timestep order.

~5.5x faster

~48 ms/step (JAX) vs. ~264 ms/step (real Fortran). Radiation excluded (simplified graybody stand-in, not real spectral transfer).

GPU: not yet measured for this full chain.

Reading this honestly

JAX-CPU is **not** a blanket win — it loses DRYCNV outright to Fortran on CPU (small-array dispatch overhead dominates at this grid size).

GPU is where JAX wins decisively on both kernels tested. That's the real case for this port.

Recommendation & Open Items

The remaining 3 modules — port or defer?

SEAICE (core thermodynamics) — port next

Physically active in every P2SAoM40-class run, not optional. Dominates polar surface energy balance — exactly where the current 0.987 global correlation check is least able to see a local error.

ATURB (PBL-top-finding) — port next

Feeds surface-flux accuracy globally, not just at the poles. Also tests whether the vmap-over-grid-cells approach holds up on harder physics.

LAKES (lkmix) — defer

Small fraction of Earth's surface — lower accuracy payoff for the effort. Revisit only if a lake-focused science need requires it.

Reporting cleanup

Drop NumPy-relative speedups from headline reporting. The target has always been **JAX vs. real Fortran** — NumPy was a control group for a narrower question, and the PBL one (`simil_numpy`) was also found to be incomplete.

Open items

- 1 Measure the full physics-chain GPU run (radiation+surface+ground together) — only the DRYCNV/PBL kernel subset has real GPU numbers today.
- 2 Fix the wind-speed convention bug in the shared FLUXES/SURFACE module files themselves — currently only worked around locally.
- 3 Port SEAICE + ATURB, then re-run the 0.987 correlation check.